

Hybrid lexical-semantic retrieval over SNOMED CT: combining two retrieval paradigms to facilitate clinical data entry

Alejandro López Osornio

Senior Implementation Support Specialist, SNOMED International

[co-authors TBD]

DRAFT, September 7, 2026

Abstract

Objectives. Searching a reference terminology such as SNOMED CT confronts two retrieval technologies of fundamentally opposite nature. Deterministic lexical matching is precise, order-independent and character-level, and lets a clinician refine a query incrementally, but it is confined to the wording and language of the curated descriptions. Learned semantic matching bridges paraphrase and language, but poorly represents abbreviations and partial, as-typed input. Their requirements and failure modes are so different that combining them naively can let the channel that is wrong for a given query degrade the one that is right. The conventional way to avoid adding a semantic channel at all has been to strengthen the lexical one (curating an additional local interface vocabulary so that more expressions match by surface form), but such vocabularies are labour-intensive to maintain and their quality degrades over time. We ask whether these two opposite paradigms can instead be made to *coexist* in one retrieval flow (each serving the input regime it is strong at, neither harming the other) over SNOMED CT’s own curated descriptions, and study the query interpretation, rank fusion and re-ranking design that makes this coexistence safe.

Materials and Methods. We implemented one reproducible instance of this architecture over the SNOMED CT International edition. It combines deterministic, order-independent multi-prefix search with three learned components: BioLORD-2023-M embeddings for semantic retrieval, an optional BGE cross-encoder for re-ranking, and a local LLM (gemma 3) for query normalization and structured clinical entity extraction. Lexical and semantic rankings are combined with Reciprocal Rank Fusion, and results can be constrained with a hierarchy (descendant) filter. As an external evaluation we measured search-only linking accuracy on 542 disease mentions from the DisTEMIST corpus (Spanish, zero-shot), resolving gold codes through SNOMED CT historical associations and

ablating query pre-processing, re-ranking, and each retrieval channel in isolation.

Results. Illustrative tests showed complementary behaviour: the lexical and semantic paths resolved different query types, while LLM pre-processing supported cross-lingual and abbreviated input. On a search-only linking evaluation over 542 disease mentions from the DisTEMIST corpus (Spanish, zero-shot), the best and computationally cheapest configuration, semantic search with re-ranking and *no* LLM query pre-processing, placed the exact concept first for 60% of mentions (accuracy@1 = 0.60) and within the top-10 candidates a picker would show for 80% (recall@10 = 0.80). Because these search-only figures are computed on gold mention spans, they isolate normalization from mention detection and are not directly comparable to the DisTEMIST shared task’s end-to-end scores. LLM query pre-processing helped dirty, lay input but slightly hurt clean terminology-like queries, and a near-miss analysis showed that many remaining cases retrieved a parent or child of the intended concept. Isolating the channels showed the two coexist without a trade-off: on this clean-normalization corpus the semantic channel carried the accuracy and adding the lexical channel neither helped nor harmed it (accuracy@1 0.60 vs 0.61), because rank fusion followed by re-ranking suppresses whichever channel is off-task; the lexical channel’s own strength, partial as-typed input, is a regime this corpus does not test. Results are specific to the evaluated configuration, not an architecture-independent estimate.

Discussion. The findings support the technical feasibility of a reference architecture that computes additional access routes over SNOMED CT’s curated descriptions, and indicate that making the lexical and semantic paradigms *coexist*, rather than choosing between them, is what makes those computed routes safe to add. This shifts rather than eliminates maintenance, from local phrase authoring toward governance and validation of retrieval services.

Conclusion. Hybrid concept retrieval lets lexical precision and semantic flexibility coexist over SNOMED CT’s maintained descriptions without a trade-off between them, offering an implementation pattern that may reduce additional local lexical curation where direct search is appropriate, while curated interface vocabularies remain useful in selected guided workflows.

Background and significance

The vocabulary challenge

Reference terminologies such as SNOMED CT are challenging to search. They are engineered for concept permanence, formal definition and computability (Cimino 1998), whereas clinicians use compact, variable and often idiosyncratic expressions. This vocabulary challenge is inherent to human communication: when people name the same thing, they choose the same word less than 20% of the time (Furnas et al. 1987).

The challenge becomes concrete at the point of *clinical data entry*, where a free-text expression must be encoded to a specific SNOMED CT concept. Traditionally this encoding is performed interactively by a clinician who types a term into a search box and selects a concept from the results, an act of iterative query refinement that depends on the user’s familiarity both with how the terminology is worded and with how the matching algorithm behaves (Rosenbloom et al. 2006). Increasingly, the same encoding is also performed automatically, by artificial-intelligence agents that extract clinical entities from narrative text and normalize each to a concept with no human in the loop. The two modes stress the terminology in different ways: an interactive user’s input varies unpredictably (from a complete term to a short, partial, as-typed fragment or abbreviation, refined keystroke by keystroke), whereas an automated agent tends to submit a complete, already-normalized phrase, often translated into a common working language. Both, however, must cross the same gap between spontaneous clinical expression and curated terminology, and a retrieval interface intended to facilitate data entry must therefore serve both the fragmentary, interactive query and the normalized, batch one over the same content.

Terminology developers and implementers have addressed this challenge through different strategies. Efficient lexical algorithms improve access to the descriptions already curated within the terminology. When that is insufficient, interface terminologies add locally curated expressions intended to anticipate how users may search. More recently, language models have made it possible to compute semantic similarity at query time instead. These approaches place the work of resolving the vocabulary mismatch, together with its corresponding maintenance burden, in different places: within the reference terminology, with the local implementer or in the retrieval system. Our hypothesis is twofold: that dynamic semantic matching can replace much of the additional local lexical curation while continuing to rely on SNOMED CT’s maintained descriptions; and that it can do so *alongside*, not instead of, algorithmic lexical search: the two coexisting so that each covers the queries the other cannot, without either degrading the other.

Curated descriptions and algorithmic lexical search

SNOMED CT already includes substantial curated lexical content: its concepts are accompanied by clinician-authored descriptions, including preferred terms and synonyms (approximately 1.6M active descriptions for approximately 368k concepts in the 2024 International edition) (SNOMED International 2024). This content reflects two important development traditions. Clinical Terms Version 3 (CTV3, formerly the Read Codes) was explicitly user-led and stressed the inclusion of the natural terms used by clinicians (O’Neil et al. 1995). In parallel, SNOMED RT was developed through a collaborative process involving the College of American Pathologists and Kaiser Permanente’s Convergent Medical Terminology (CMT) team (Levy et al. 1998). The merger of CTV3 and SNOMED RT brought these traditions together in SNOMED CT. This is also

the genealogy of the clinician-facing vocabulary embedded in SNOMED CT: a centrally curated interface resource distributed and maintained with the reference terminology, rather than separately by each implementation.

In routine use, modern terminology servers combine this lexical content with fast, responsive search. Part-word retrieval has a long lineage in UK clinical terminology systems: clinicians searched Read Codes and CTV3 through the NHS Clinical Terminology Browser by entering one or more words or part-words (Brown et al. 2003), and later tooling for Read-coded data formalized searches using multiple word-stubs occurring anywhere in a description (Olier et al. 2016). The *multi-prefix, any-order* technique used here follows this interaction pattern: users enter only the first characters of one or more words, and each fragment is matched against the beginning of a word in the description, regardless of word order. Thus, the fragments *myo inf* can retrieve *myocardial infarction*, and reversing them does not change the match. This tolerates incomplete words and natural variation in word order while preserving precision by requiring every entered prefix. Because inflected and derived forms often share their initial characters, it also provides lightweight morphological normalization across variations such as number and, in languages that mark it, gender. This benefit is language-dependent and does not amount to full lemmatization. Clinicians can start with only a few characters and refine them in response to the results until they reach the concept they intend. This simple interaction makes the user an active participant in resolving the vocabulary mismatch, without requiring every possible expression to be anticipated in advance.

Additional local lexical curation

Some implementations, however, need to reduce that dependence on individual search skills. Adoption constraints, limited training opportunities or the demands of a particular workflow may call for more natural queries and more guided concept selection. The conventional response is an additional *interface terminology*: a “systematic collection of health care-related phrases that support clinicians’ entry of patient-related information” (Rosenbloom et al. 2006). Added to the descriptions distributed with SNOMED CT, these local vocabularies map further colloquial expressions, abbreviations and preferred terms to concepts in the reference terminology. In effect, an interface terminology strengthens the lexical channel: it supplies more surface forms for each concept so that more queries succeed by wording alone. Until the recent maturation of efficient semantic retrieval, this was essentially the only way to broaden coverage without depending on the user’s search skill: adding a second, semantic channel simply was not a practical option. Kaiser Permanente’s CMT is a prominent example: it combined SNOMED CT with clinician-facing terms, mappings and context-specific subsets as an enterprise terminology resource (Dolin et al. 2004).

Published operational experience shows that this convenience is sustained by ongoing expert work. At the University of Nebraska Medical Center, a just-in-

time process allowed clinicians to submit problems that were not yet represented in its 12,000-term lexicon. Evaluating an unlisted problem, finding an appropriate code, notifying the clinician and coding the record required an average of 30 minutes of expert coder time; unresolved terms were escalated to a vocabulary team (Warren et al. 1998). At Hospital Italiano de Buenos Aires, a terminology service developed through continuous user feedback recognized 140 ways to describe arterial hypertension and thousands of Spanish clinical expressions absent from a general-language dictionary. Its permissive policy improved user acceptance, but 67% of concepts in the problem-list subset required post-coordination; local mappings required additional manual work; and rules based on SNOMED CT hierarchies had to be reviewed with each new release (González Bernaldo de Quirós et al. 2018). At enterprise scale, early CMT tooling was designed to reconcile conflicting local enhancements developed by Kaiser Permanente and Mayo Clinic and to distribute updates compatible with those extensions (Campbell et al. 1996). These are successful systems, but their success depends on specialized teams, governance and tooling.

That work creates a substantial maintenance obligation for each implementation. Every additional interface term must be authored, mapped to a concept and kept aligned as the reference terminology evolves. More generally, terminology maintenance has been characterized as complex, resource-intensive and time-consuming, with increasing size requiring more extensive change management, validation, consistency checking, versioning and editing support (Bakhshi-Raiez et al. 2008). SNOMED International notes that a custom interface terminology is “challenging to maintain and keep up to date with new SNOMED versions, with risks of quality issues such as duplication and ambiguity,” and recommends, where feasible, using SNOMED CT’s own descriptions as the interface terminology (SNOMED International 2024). Without sustained maintenance, detailed local interface terminologies and the value sets and mappings that depend on them can fall out of alignment with evolving reference terminologies, introducing duplication, ambiguity and inconsistencies across versions (Bakhshi-Raiez et al. 2008; SNOMED International 2024; González Bernaldo de Quirós et al. 2018). The tension between flexible, clinician-friendly expression and structured, maintainable representation is long-recognized (Rosenbloom et al. 2011). However, the imbalance between unbounded expressive variety and finite local curation capacity makes an exhaustively curated additional interface vocabulary difficult to sustain.

This frames the question the rest of the paper pursues. If the additional curation could be replaced by retrieval *computed* at query time, the maintenance burden would fall, but only if that computed route can run *beside* the terminology’s own lexical search without degrading it, because otherwise one has merely traded a maintenance problem for an accuracy problem. Reducing local curation and making two retrieval paradigms of opposite nature coexist are therefore the same problem seen from two sides: the motivation and the technical crux.

Dynamic semantic matching

Recent language technologies offer another way to facilitate search without creating a separate local vocabulary. Biomedical text-representation models place synonyms and paraphrases close in a shared semantic space even without lexical overlap (Remy et al. 2024; Liu et al. 2021); this approach has powered semantic search over large clinical ontologies (notably the CSIRO group’s learned semantic search over SNOMED CT, which outperformed a lexical baseline (Ngo et al. 2021)) and unsupervised concept annotation against SNOMED CT (Abdulnazar et al. 2024). Crucially, semantic matching does not replace the lexical curation performed by SNOMED CT: embeddings are generated from its existing descriptions, so the terminology’s curated lexical content anchors the semantic representation. Combined with LLMs that can normalize and translate clinician shorthand, these methods can dynamically bridge the gap between a natural query and those descriptions, potentially replacing much of the manual expansion with additional local synonyms.

Two retrieval technologies of opposite nature

These two ways of reaching a concept are of opposite nature in almost every respect. Lexical matching is deterministic, character-level and precise: it does exactly what the user’s keystrokes specify, tolerates partial and out-of-order input, and lets the clinician steer the search interactively, but it cannot cross a wording or language gap for which it was given no shared tokens. Semantic matching is learned, meaning-level and approximate: it crosses paraphrase and language with ease, but it poorly represents abbreviations and half-typed fragments, and cannot be steered keystroke by keystroke. The two are strong and weak on disjoint inputs, and their scores are not even on the same scale. Combining them is therefore not obviously beneficial: for any given query one channel is usually off-task, and a naive combination lets that off-task channel inject distractors that displace the concept the on-task channel has already found. Nor can the difficulty be sidestepped by routing each query to the method that suits it, because in real use the boundary is not stable even within a single user: the same clinician may type a complete term one moment and a two-letter prefix the next, so which regime a query belongs to cannot be read reliably from the query itself, let alone predicted in advance. Delegating that routing decision to a language model on every query (or, in an incremental search box, on every keystroke) would add latency and computational cost out of proportion to the retrieval it is meant to steer. The central design question of this paper is whether two technologies this different can be made to *coexist* in one retrieval flow (each contributing where it is strong, neither degrading the other) without requiring the system to know in advance which regime a query belongs to.

Objective

We propose a generalizable reference architecture for retrieving SNOMED CT concepts from free-text clinical expressions, and illustrate it with a locally runnable, open-source implementation and a preliminary external evaluation on a Spanish gold-standard corpus. Its organizing question is how retrieval technologies of opposite nature, deterministic lexical matching and learned semantic matching, can coexist over the same curated descriptions without degrading one another. The architecture separates query interpretation, complementary lexical and semantic candidate generation, rank fusion, optional re-ranking, and terminology-aware constraints. Our aim is not to prescribe a particular technology stack or introduce a new state-of-the-art entity linker. Instead, we examine *hybrid concept retrieval* as an alternative implementation pattern between direct lexical search and an additional hand-maintained interface vocabulary. Publishing the reference implementation makes its design choices inspectable and allows others to adapt and evaluate the pattern in their own settings.

Materials and methods

Reference architecture

At an abstract level, the proposed architecture comprises five separable functions: (1) interpreting or normalizing the input query; (2) generating candidates through complementary lexical and semantic retrieval paths; (3) combining evidence from rankings whose scores may not be directly comparable; (4) optionally re-ranking a small candidate set with a more computationally intensive model; and (5) applying terminology-aware constraints and deterministic safeguards. The architecture does not depend on a particular LLM, embedding model, re-ranker, database or fusion algorithm. Implementers can substitute or omit components according to their languages, latency requirements, available infrastructure and governance constraints. Its central principle is to preserve the precision and responsiveness of lexical matching while adding a semantic route for expressions that do not share the terminology’s wording; and, critically, to combine the two so that neither degrades the other. Functions (3) and (4), fusion and re-ranking, are precisely where that coexistence is secured: fusion admits candidates from whichever channel is on-task, and re-ranking demotes the distractors contributed by the channel that is off-task for a given query. Figure 1 summarizes these functions and the flow between them.

The architecture is also language-independent by construction. Its two language-sensitive functions, query interpretation and semantic retrieval, can be instantiated with multilingual models, which already exist off the shelf (multilingual biomedical bi-encoders, LLMs and cross-encoders), so a query in one language is resolved against descriptions in another without any per-language curation: the lexical path stays tied to the language of the indexed descriptions, while the semantic path carries meaning across languages. Cross-lingual retrieval is there-




Figure 1: Generalizable reference architecture for hybrid SNOMED CT concept retrieval. Dashed borders indicate optional functions. Components within each function may be substituted according to local requirements; arrows show the flow of queries, candidates and rankings.

fore an available design choice rather than a per-deployment research problem.

Reference implementation

The system evaluated here is one concrete instance of this model. A local LLM normalizes short queries or extracts clinical entities from longer text; deterministic multi-prefix matching supplies the lexical candidates; a biomedical bi-encoder supplies the semantic candidates; Reciprocal Rank Fusion combines both rankings; an optional cross-encoder re-ranks the fused candidates; and hierarchy filters constrain results to an appropriate SNOMED CT domain. In this implementation, retrieval, fusion and hierarchy filtering are served from a single PostgreSQL instance and executed in one query plan. The following sections describe these choices to make the example reproducible, rather than as requirements of the general architecture. The open-source implementation also provides an executable reference in which components can be inspected, replaced and re-evaluated against different languages, SNOMED CT editions, corpora and infrastructure.

Lexical channel

Active descriptions are indexed as PostgreSQL full-text vectors with a custom accent- and case-insensitive, non-stemming configuration and a GIN index. At query time, each fragment entered by the user becomes a prefix clause and the clauses are joined by AND; every fragment must therefore match the initial characters of a word in the description, in any order. For example, *myo inf* is represented as `myo:* & inf:*` and matches *Myocardial infarction* as readily as *inf myo*. Shared prefixes also absorb some language-dependent inflectional and derivational variation without explicit stemming or lemmatization. This supports incremental search with minimal typing while avoiding the broad result sets produced by matching only some of the entered prefixes. Ranking uses `ts_rank` with length and unique-word normalization so that concise canonical descriptions outrank verbose ones (PostgreSQL Global Development Group 2026).

Semantic channel

Each description is embedded with BioLORD-2023-M, a biomedical, multilingual sentence bi-encoder (Remy et al. 2024) whose lineage traces to self-aligned biomedical representations (Liu et al. 2021), and stored in a `pgvector` column with an HNSW index for approximate nearest-neighbour search by cosine similarity.

Fusion and re-ranking

The two rankings are combined with Reciprocal Rank Fusion (RRF) (Cormack et al. 2009), which requires no score calibration between channels. An optional cross-encoder re-ranker (BGE-reranker-v2-m3, built on a multilingual encoder

(Chen et al. 2024)) refines the fused top- k ; its score is min-max blended with the RRF score rather than allowed to override it. A deterministic rule pins exact-match descriptions to the top.

Hierarchy filter

A precomputed transitive closure of the IS-A graph allows any search to be constrained to the descendants-or-self of a concept (e.g. only *Clinical findings*), turning “is concept C under X?” into a single indexed array-containment lookup. This matters because a single surface form often has near-neighbours in several SNOMED CT hierarchies (the same words can name a *finding*, a *morphologic abnormality*, a *body structure* or a *substance*), and constraining the search to the intended domain removes those cross-hierarchy homonyms before they compete for the top rank. In the entity-extraction path the constraint is supplied automatically by the extracted semantic type, so the LLM’s typing decision feeds directly into retrieval precision.

LLM roles

A local LLM (gemma 3 (Gemma Team, Google DeepMind 2025), served through an OpenAI-compatible endpoint) plays two roles. (1) *Query normalization*: it translates a short clinician query from any language to one canonical English clinical term and expands abbreviations, so that both retrieval channels operate on clean English; the embedding model’s cross-lingual nature also allows the system to degrade gracefully without the LLM. (2) *Entity extraction*: given a free-text note, it returns structured entities as JSON, including the verbatim span, semantic type, assertion/temporal context (present/absent/past/unknown), laterality, severity and an English `clinicalTerm`. Each extracted entity is then mapped independently through the hybrid engine, using its normalized `clinicalTerm` as the query and its semantic type as the hierarchy filter (falling back to an unfiltered search when the type filter would exclude everything).

Data and implementation

We used the RF2 Snapshot view of the 1 June 2026 SNOMED CT International Edition production release. We loaded the active English-language descriptions of active concepts: 1,024,825 descriptions (381,856 fully specified names and 642,969 synonyms) for 381,856 concepts; of these, 1,020,940 unique normalized description strings were embedded. A transitive closure of the IS-A relationships was precomputed over all 381,856 active concepts to support the hierarchy (descendant) filter. The stack runs locally on a single workstation (Apple Silicon), with the database in a container and all model inference in-process; no clinical data leaves the machine. The full pipeline, indexes and demonstration interface are reproducible from a public repository [TODO: URL].

External evaluation

Choice of corpus

Openly licensed gold standards that link *free clinical text* to SNOMED CT are scarce. Most English clinical-note resources that normalize to SNOMED CT or the UMLS require a data-use agreement, and general clinical NLP corpora do not target SNOMED CT. The Spanish corpora produced by the Barcelona Supercomputing Center for the BioASQ/BioCreative shared tasks are a notable exception: they are released under CC BY 4.0 and every mention is normalized to SNOMED CT. We evaluated on **DisTEMIST** (BioASQ 2022), 1,000 Spanish clinical case reports with disease mentions linked to SNOMED CT (Miranda-Escalada et al. 2022), because diseases map cleanly onto the *Clinical finding* hierarchy. Two caveats follow from this choice and we treat them explicitly rather than hide them. First, the corpus is **Spanish**; for a multilingual system this is less a limitation than a cross-lingual stress test, since the query and the terminology descriptions are in different languages. Second, the texts are drawn from **published case reports**, whose language is more edited than bedside notes; the evaluation therefore probes concept normalization more than the full messiness of real electronic health record text. We use this corpus because it is among the few open, SNOMED-CT-linked options, and read the results with these limits in mind.

Task and configurations

To measure the *search* capability independently of our entity extractor, we ran a **search-only** evaluation: each gold mention’s verbatim span is submitted as a query and we record the rank of the gold SNOMED CT code; this is the corpus’s own entity-linking (normalization) subtask. We swept the two learned, optional components of the search pipeline as a 2×2 design, LLM query pre-processing (gemma normalization/translation) on/off × cross-encoder re-ranking on/off, and report three of the four cells. The lexical and semantic channels and their RRF fusion were always active; the hierarchy filter was fixed to *Clinical finding* (the corpus is diseases); and we retrieved $k = 10$ candidates.

Edition-drift resolution

The gold was annotated against a circa-2022 edition; by our loaded release (International 2026-06-01) some gold concepts had been inactivated. Because a fixed gold standard drifts as the terminology evolves (itself an instance of the maintenance problem this paper is about), we resolved every gold code through SNOMED CT’s **historical-association** reference set (in priority order SAME AS → REPLACED BY → POSSIBLY REPLACED BY → POSSIBLY EQUIVALENT TO → ALTERNATIVE) to its current active concept before scoring.

Scope

The system targets *disorder* and *finding* concepts. Where the annotators coded a mention to a *substance* or *morphologic-abnormality* concept for the same clinical entity (for example, *renal stone* as a substance rather than as a disorder), a system that returns the disorder can never match by identifier. We therefore scored only gold whose resolved concept is a disorder or finding, so that representation choices outside the system’s target do not create artificial misses.

Metrics

We chose metrics that mirror how a clinician uses a concept picker. In a typical search interface the user types a term and selects from a short ranked list, so what matters is whether the intended concept appears *near the top*. We therefore report, under a **strict** rule (our concept identifier equals the resolved gold): **accuracy@1**: the top candidate is exactly the gold concept; **recall@5** and **recall@10**: the gold concept is among the first 5 or 10 candidates, i.e. the options a user would see in a drop-down list without, or with a small scroll; and **mean reciprocal rank (MRR)**, which rewards a higher position of the gold concept. We deliberately foreground recall@5 and recall@10 because they answer the practical question for a picker interface, “would the right concept be on the screen for the user to choose?”, rather than requiring the system to be right on its single best guess. Separately, because a *parent* or *child* of the intended concept can be an acceptable or navigable alternative, we report a **near-miss** rate: the additional fraction of mentions (beyond the exact matches) whose top-k list contains an ancestor or descendant of the gold. The near-miss rate is thus an increment on strict recall, not a replacement for it: the two together sum to a “same-lineage” recall. It is an upper bound on “clinically acceptable” rather than a hard accuracy, since a hierarchical neighbour is not always an acceptable substitute; confirming that would require clinician adjudication.

Results

Feasibility of the reference architecture

The open-source reference implementation instantiated all five functions of the proposed architecture in an end-to-end pipeline. For short expressions, it generated lexical and semantic candidates and returned a fused ranking. For longer notes, it first extracted and normalized clinical entities and then retrieved candidates independently for each entity. Optional re-ranking and hierarchy constraints operated on the same candidate flow, allowing the architecture to support both interactive concept search and note-level entity linking. The same open-source application, an interactive search interface, a note-level entity-extraction view and a batch evaluation harness, was the instrument for both the qualitative observations and the quantitative results that follow. This establishes the technical feasibility of the architectural pattern; the observations

and metrics below characterize this particular implementation.

Illustrative component behaviour

The examples exposed complementary roles for the components rather than a consistent advantage for any single model. The semantic path recovered lay phrasing with no shared tokens (“water on the knee” → *Effusion of knee joint*), whereas deterministic multi-prefix matching retrieved a clinician abbreviation (“inf myo” → *Myocardial infarction*) that was represented poorly by the embedding model. Query interpretation and cross-lingual retrieval supported inputs such as “radiografía de tórax” → *Plain X-ray of chest* and “EPOC” → *Chronic obstructive pulmonary disease*. For longer notes, the LLM extracted entities and represented assertion separately from concept identity, mapping “no fever” to *Fever* with an absent flag and “history of stroke” to *Cerebrovascular accident* with a past flag. In a Spanish note, entities were extracted verbatim, normalized to English `clinicalTerms`, and mapped to the expected SNOMED CT concepts (e.g. *diabetes mellitus tipo 2* → *Diabetes mellitus type 2*; *disnea de esfuerzo* → *Dyspnea on exertion*).

External evaluation on DisTEMIST

We evaluated the first 100 DisTEMIST training documents, which contain 588 unique disease mentions; 20 gold codes were recovered through historical associations and 46 mentions were scoped out as non-disorder/finding, leaving **542** mentions. Table 1 reports the three search configurations.

Table 1: Search-only linking on 542 DisTEMIST disease mentions (Spanish, zero-shot). *PP* = LLM query pre-processing (normalization); *RR* = cross-encoder re-ranking. *MRR* = mean reciprocal rank of the gold concept. The strict columns *acc@1*, *R@5* and *R@10* (accuracy@1, recall@5, recall@10) require the exact concept identifier and read as a picker list: is the gold concept on the screen the user would see? *+near@10* is the additional fraction whose top-10 held an ancestor or descendant of the gold (the near-miss increment), and *=lin@10* is the two summed (same-lineage@10): the most optimistic reading, in which any hierarchical neighbour counts as acceptable. Best value per column in **bold**.

PP	RR	MRR	acc@1	R@5	R@10	+near@10	=lin@10
off	on	0.67	0.60	0.76	0.80	+0.09	0.89
on	off	0.60	0.54	0.70	0.75	+0.12	0.87
on	on	0.61	0.54	0.69	0.75	+0.12	0.87

Read the numbers as a concept picker: recall@5 and recall@10 are the chance that the intended concept is on the short list a clinician would scan and select from, while accuracy@1 is the chance the system’s first candidate is already correct. Three findings stand out. First, the **cheapest configuration is also the best**: with query pre-processing *off* and re-ranking *on*, the exact concept is

the top candidate for 60% of mentions ($\text{accuracy@1} = 0.60$) and appears within the top-10 picker list for 80% ($\text{recall@10} = 0.80$); for a further 9% of mentions a parent or child of the intended concept is in that list ($\text{near-miss@10} = +0.09$). Because this configuration makes no per-query LLM call, it is simultaneously the most accurate and the least expensive to run. Second, **LLM query pre-processing hurts on this task**: enabling it lowers accuracy@1 from 0.60 to 0.54. The gold mentions are already clean, terminology-like disease terms, and the multilingual embedding maps them cross-lingually without help; the LLM’s normalization to English instead shifts the semantic neighbourhood (for example, *infertilidad* $\rightarrow$ *infertility* pulls the specific “X infertility” disorders above the exact generic *Infertile*). The LLM’s value is for the dirty, lay or abbreviated input of interactive use, not for well-formed terminology strings. Third, **re-ranking helps** when it is not preceded by aggressive normalization, improving both accuracy@1 and MRR; the two components are genuinely optional and the optimal configuration is task-dependent.

The **MRR of 0.67** for the best configuration is worth reading carefully, because it characterizes the ordering rather than mere presence. MRR is a mean of reciprocal ranks in which unretrieved mentions score zero, so it cannot simply be inverted into an “average rank.” Its value here is only 0.07 above accuracy@1 (0.60), and that small gap is the informative part: of the mentions the system does retrieve, almost all sit at or immediately below the top, and very few correct concepts are buried deep in the list. Decomposing the 542 mentions, roughly 60% are exact at rank 1, about 20% are retrieved but at a lower rank (contributing an average reciprocal rank near 0.35, i.e. around rank 3), and the remaining 20% are not retrieved within the disorder/finding scope at all. The practical implication is that the residual difficulty is one of *recall* (the concepts the system never surfaces) rather than of *ranking*: when the intended concept is present, the picker rarely makes the clinician scroll for it. Read together, the columns bound the outcome from both sides: strict $\text{accuracy@1} = 0.60$ is the pessimistic floor (the first candidate is exactly right), while $\text{same-lineage@10} = 0.89$ is the optimistic ceiling (the top-10 holds the exact concept or a hierarchical neighbour a clinician might accept), with the true operational value depending on how a given workflow treats near-misses.

These are **zero-shot** results (no training or fine-tuning on the corpus). They are not directly comparable to the results reported for the DisTEMIST shared task: its best supervised linking system reached a micro-averaged F1 of 0.57 (Miranda-Escalada et al. 2022), but that figure is *end-to-end*, scoring mention detection and normalization together, whereas our search-only accuracy is computed on gold mention spans and isolates the normalization step that is this paper’s subject. The two are of similar magnitude but measure different tasks; a like-for-like comparison would require scoring our full extraction-and-mapping pipeline end-to-end, which (as the present study concerns the search step rather than entity extraction) we leave to future work. The near-miss rate shows that a meaningful share of the strict “errors” are hierarchically adjacent rather than wrong: at rank 1, beyond the 0.60 exact matches, a further 0.15 of mentions

return a parent or child concept as the top candidate. Whether such a neighbour is acceptable is a clinical judgement, so we treat this as an upper bound rather than as accuracy. A further residual class, credited by neither the exact nor the near-miss count, is the annotators’ choice of a *substance* or *morphologic-abnormality* concept where the system returns the corresponding *disorder*, a genuine representation difference in the gold itself, not a retrieval failure.

Running the full extraction-then-mapping pipeline over the same notes (rather than the gold spans) lowers these figures, which isolates the contribution of the extraction stage: the LLM occasionally mis-types or omits an entity, and long notes must be chunked to prevent the small model from degenerating and dropping their tail. In this implementation the search is the stronger component; extraction is where accuracy is currently lost.

Coexistence of the two retrieval channels

To test whether the two channels coexist without a trade-off, we repeated the best serving configuration (no query pre-processing, re-ranking on) while restricting retrieval to a single channel (Table 2). The semantic channel accounts for essentially all of the accuracy on this corpus, and fusing the lexical channel back in leaves it unchanged (accuracy@1 0.60 vs 0.61; the difference is a single mention). The lexical channel alone collapses (accuracy@1 0.07) for a specific reason: the queries are Spanish while the indexed descriptions are English, so without the query-normalization step the lexical channel has almost no tokens to match, whereas the multilingual embedding bridges the two languages directly. The informative result is the near-identity of the semantic and fused rows: on clean cross-lingual normalization, adding the lexical channel neither helps nor harms.

Table 2: Retrieval-channel ablation on the same 542 DisTEMIST mentions, in the best serving configuration (no query pre-processing, re-ranking on); only the retrieval channel varies. *both* is the Reciprocal Rank Fusion of the two channels.

channel	acc@1	recall@10	MRR
lexical only	0.07	0.07	0.07
semantic only	0.61	0.80	0.68
lexical + semantic (fused)	0.60	0.80	0.67

Contribution of the hierarchy filter

The hierarchy filter, fixed to *Clinical finding* throughout, contributes at the top of the ranking. Repeating the best configuration without it lowers accuracy@1 from 0.60 to 0.56 and MRR from 0.67 to 0.64, while recall@10 barely moves (0.80 to 0.78). The pattern is that of a precision tool: the correct concept is usually retrieved within the top-10 either way, but without the domain

constraint, concepts from other SNOMED CT hierarchies (a *substance*, a *morphologic abnormality* or a *body structure* sharing the query’s words) sit above it and displace it from rank 1. Because the constraint is a subtree test over a precomputed transitive closure, this precision gain costs a single indexed lookup.

Discussion

Principal findings. This study contributes a generalizable reference architecture and an executable, open-source implementation. The implementation demonstrates the technical feasibility of combining query interpretation, complementary lexical and semantic retrieval, rank fusion, optional re-ranking and terminology-aware constraints in a single concept retrieval flow. Illustrative examples showed that these functions address different forms of vocabulary mismatch, while the preliminary external evaluation established a reproducible baseline for one concrete configuration. A more specific finding sharpens the architectural claim: two retrieval technologies of opposite nature can be fused without a trade-off, provided re-ranking suppresses whichever channel is off-task for a given query (developed below). These findings support the coherence and feasibility of the architectural pattern. They do not establish that the selected components are optimal, that the observed performance transfers to other implementations, or that the system is ready for clinical use.

Implications for terminology infrastructure. Hybrid concept retrieval provides an intermediate strategy between expecting clinicians to resolve every mismatch through lexical query refinement and attempting to anticipate every likely expression in an additional interface vocabulary. A responsive lexical path preserves familiar, precise interaction with SNOMED CT descriptions, while a semantic path computes further candidates when the query uses different wording. Both paths remain anchored in lexical content curated and distributed with SNOMED CT. This may reduce local term-by-term authoring and mapping, but it does not make maintenance disappear. Instead, it shifts work toward governing models, indexes, constraints, terminology releases and regression tests. Some of this work can be automated and shared, whereas exhaustive local phrase curation grows with the variety of clinician language (Furnas et al. 1987). When the retrieval models are multilingual, as in our implementation, that computed bridge spans languages as well, so a deployment need not curate colloquial expressions language by language. Additional curated interfaces remain appropriate when highly guided workflows, local policy or limited opportunities for user training make them valuable. Using SNOMED CT’s own descriptions directly where feasible also follows SNOMED International guidance (SNOMED International 2024).

Coexistence without a trade-off. The channel ablation (Table 2) shows the semantic channel carries clean disease normalization on DisTEMIST and that fusing the lexical channel back in costs nothing measurable. Read naively this argues for dropping the lexical path, but that reading is an artifact of the

benchmark: DisTEMIST mentions are complete, edited disease terms, so they never exercise what the lexical channel is uniquely good at: order-independent matching of partial, as-typed input such as the abbreviation “inf myo” for *Myocardial infarction*, a fragment the embedding model represents poorly (Section *Illustrative component behaviour*). The two channels serve different regimes (whole-word, paraphrase and cross-lingual matching for the semantic path; prefix, abbreviation and exact-fragment matching for the lexical path), and a single query does not announce which regime it belongs to. Rank fusion followed by cross-encoder re-ranking is what lets one pipeline keep both worlds: fusion contributes candidates from whichever channel is on-task, and the re-ranker suppresses the off-task channel so it does not degrade the result. This matters because naive equal-weight fusion *without* re-ranking does lose ground to the stronger single channel (the off-task lexical path injects distractors that displace the correct concept), whereas re-ranking recovers the loss and returns fusion to parity with the best channel. The practical payoff is regime-robustness without query routing: the same system serves a Spanish disease term and a half-typed English abbreviation, without having to classify the query in advance to choose a channel. The benchmark measures only one of these regimes, so it should be read as evidence that hybridization is *safe* for clean normalization, not as evidence that the lexical channel is dispensable.

These two regimes map onto the two modes of clinical data entry noted at the outset, and the mapping suggests that the value of hybridization is not uniform across callers. An automated agent typically submits a complete, already-normalized term (precisely the input the semantic channel handles best), so for agent-driven encoding the semantic path alone may be sufficient, and the lexical channel adds little. A clinician entering data interactively types partial, abbreviated, as-typed fragments and refines them keystroke by keystroke: the regime in which the lexical channel is indispensable and in which the hybrid most improves usability. Hybridization therefore earns its keep chiefly at the human, interactive end of the spectrum, while a well-formed agent query may be served by semantics alone, though the same pipeline serves both without having to know which is calling.

The hierarchy filter as a cheap precision lever. Fusing the two channels is not the only place the architecture exploits SNOMED CT; constraining a search to the expected hierarchy raised top-rank precision at negligible cost (accuracy@1 +0.05, MRR +0.03 in our evaluation) by removing homonyms drawn from unintended hierarchies. This is where the terminology’s *structure*, not only its curated descriptions, is put to work, and where an automated agent’s (or the LLM extractor’s) semantic-type classification pays off directly: a correct type narrows retrieval to the right domain for the price of one indexed subtree test. The same mechanism carries a recall risk, since a wrong type can exclude the intended concept, which the implementation mitigates by falling back to an unconstrained search when the filtered query returns nothing. For the agent-driven encoding discussed above, this makes entity typing not an incidental preprocessing step but a lever on retrieval precision.

Architecture, implementation and validation. The five architectural functions define separable responsibilities rather than fixed products. Implementations may use different language models, embedding models, lexical algorithms, fusion methods, databases and constraints, and their results may therefore differ substantially. Performance claims should attach to a specified, versioned implementation, while support for the reference architecture should accumulate across implementations and settings. The open-source system and evaluation materials provide a starting point for a reusable evaluation harness: canonical queries, external corpora, ranking metrics, latency measures and failure cases can be rerun when a component changes. Automated regression testing can make iterative improvement less costly, but it cannot replace validation against the intended users and clinical use case. The required depth of that validation should reflect the consequences of an incorrect match and whether the output is advisory, user-confirmed or recorded without review.

Relation to prior work. Semantic search over clinical ontologies and SapBERT/BioLORD-style bi-encoders established that embeddings can recover concepts under vocabulary mismatch (Ngo et al. 2021; Liu et al. 2021; Remy et al. 2024; Abdulnazar et al. 2024), while recent scoping work catalogues uses of LLMs with SNOMED CT (VERIFY authors 2024). The distinctive contribution here is architectural rather than a new retrieval model. The closest precedent, the CSIRO system, benchmarked a learned semantic model against a lexical baseline and found the semantic model superior (Ngo et al. 2021); our channel ablation replicates that ordering on a clean-normalization corpus. Where that line of work asks *which single method retrieves best*, we ask *how* a deterministic lexical method and a learned semantic one can coexist in one flow, so that the strength of each is available without the channel that is wrong for a given query degrading the result. The design assigns bounded roles to an LLM for query interpretation and entity extraction, preserves a deterministic lexical path, adds semantic candidate generation, combines rankings without requiring score equivalence, permits more expensive re-ranking only over a small candidate set, and applies constraints derived from the terminology. It connects these technical choices to the operational question of whether access to SNOMED CT can be made more forgiving without creating another large vocabulary to maintain.

Limitations. The external evaluation covers one corpus, one language, one concept type (diseases) and the search-only linking task on gold mention spans; end-to-end extraction, other concept types (procedures, findings), other languages, and real electronic-health-record notes rather than published case reports remain to be tested. The corpus being Spanish and drawn from case reports means the input is cleaner than bedside documentation. The strict/hierarchy-aware gap indicates that the exact-identifier metric understates semantically acceptable retrieval, but neither metric was confirmed by blinded clinical adjudication. We ablated the two optional components (query pre-processing and re-ranking), the hierarchy filter, and the two retrieval channels in isolation, but did not substitute the embedding model, lexical algorithm or fusion method,

so the contribution of those functions is not estimated. We also did not compare clinician search performance, measure longitudinal maintenance effort, or evaluate safety in a production workflow. LLM entity typing remains imperfect, extraction latency is LLM-bound, and mapping to a single best concept omits post-coordination. Open source makes the design choices transparent and facilitates independent evaluation, but does not remove the need to validate each implementation against its intended users, terminology edition, language and use case. Next steps include larger evaluations across findings, diseases and procedures; component substitution and ablation experiments; hierarchy-aware and clinician-adjudicated metrics; usability studies; an evaluation on interactive, partially-typed input that exercises the lexical channel’s regime, which a corpus of complete disease terms does not; learned or weighted rank fusion, since equal-weight fusion proved sub-optimal when one channel is off-task; and longitudinal measurement of maintenance work.

Conclusion

SNOMED CT’s curated lexical content can serve as the foundation for an interface that is partly *computed* at query time. The proposed reference architecture offers a testable, implementation-independent path for combining lexical precision with semantic flexibility so that two technologies of opposite nature coexist over the same descriptions without degrading each other, a coexistence secured by rank fusion and re-ranking rather than by choosing between the paradigms. Its open-source implementation demonstrates feasibility and enables adaptation and re-evaluation, but each resulting system requires validation in its intended setting. The benefit of that coexistence appears to fall unevenly on the two forms of clinical data entry: an automated agent, which submits well-formed normalized terms, may be served by the semantic path alone, whereas an interactive human, who types partial and abbreviated fragments, gains the most usability from the hybrid. This approach may reduce additional local lexical curation, while hand-built interface terminologies retain a role in guided, low-friction clinical workflows.

Acknowledgements

AI coding and writing assistants were used during this work. OpenAI Codex and Claude Code (Anthropic) were both used during manuscript preparation to assist with language editing, structural revision and drafting; Claude Code was additionally used as a programming assistant to help develop the reference implementation, tooling and evaluation frameworks. All AI-assisted content and code was critically reviewed, verified and revised by the authors, who take full responsibility for the final manuscript.

Funding

This work was funded by SNOMED International.

Competing interests

The author is employed by SNOMED International, the organization that develops and maintains SNOMED CT, the terminology studied in this work. SNOMED International funded the work. The author declares no other competing interests.

Data and code availability

The reference implementation, configuration and evaluation materials are available as open-source software [TODO: repository URL]. They are provided to make the architectural choices inspectable and to support adaptation and re-evaluation in other environments. DisTEMIST (Miranda-Escalada et al. 2022) is available under CC BY 4.0 from its authors and is not redistributed here.

References

- Abdulnazar, Akhila, Roland Roller, Stefan Schulz, and Markus Kreuzthaler. 2024. “Unsupervised SapBERT-Based Bi-Encoders for Medical Concept Annotation of Clinical Narratives with SNOMED CT.” *Digital Health* 10. <https://doi.org/10.1177/20552076241288681>.
- Bakhshi-Raiez, Ferishta, Ronald Cornet, and Nicolette F. de Keizer. 2008. “Development and Application of a Framework for Maintenance of Medical Terminological Systems.” *Journal of the American Medical Informatics Association* 15 (5): 687–700. <https://doi.org/10.1197/jamia.M2531>.
- Brown, Philip J. B., Victoria Warmington, Michael Laurence, and A. Toby Prevost. 2003. “Randomised Crossover Trial Comparing the Performance of Clinical Terms Version 3 and Read Codes 5 Byte Set Coding Schemes in General Practice.” *BMJ* 326 (7399): 1127. <https://doi.org/10.1136/bmj.326.7399.1127>.
- Campbell, Keith E., Simon P. Cohn, Christopher G. Chute, Glenn D. Rennels, and Edward H. Shortliffe. 1996. “Gálapagos: Computer-Based Support for Evolution of a Convergent Medical Terminology.” *Proceedings of the AMIA Annual Fall Symposium*, 269–73.
- Chen, Jianlv, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. 2024. “BGE M3-Embedding: Multi-Lingual, Multi-Functionality, Multi-

Granularity Text Embeddings Through Self-Knowledge Distillation.” *arXiv Preprint arXiv:2402.03216*.

- Cimino, James J. 1998. “Desiderata for Controlled Medical Vocabularies in the Twenty-First Century.” *Methods of Information in Medicine* 37 (4-5): 394–403.
- Cormack, Gordon V., Charles L. A. Clarke, and Stefan Büttcher. 2009. “Reciprocal Rank Fusion Outperforms Condorcet and Individual Rank Learning Methods.” *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 758–59. <https://doi.org/10.1145/1571941.1572114>.
- Dolin, Robert H., John E. Mattison, Simon P. Cohn, Keith E. Campbell, et al. 2004. “Kaiser Permanente’s Convergent Medical Terminology.” *Studies in Health Technology and Informatics* 107: 346–50. <https://doi.org/10.3233/978-1-60750-949-3-346>.
- Furnas, G. W., T. K. Landauer, L. M. Gomez, and S. T. Dumais. 1987. “The Vocabulary Problem in Human-System Communication.” *Communications of the ACM* 30 (11): 964–71. <https://doi.org/10.1145/32206.32212>.
- Gemma Team, Google DeepMind. 2025. *Gemma 3 Technical Report*. Google DeepMind.
- González Bernaldo de Quirós, Fernán, Carlos Otero, and Daniel Luna. 2018. “Terminology Services: Standard Terminologies to Control Health Vocabulary: Experience at the Hospital Italiano de Buenos Aires.” *Yearbook of Medical Informatics* 27 (1): 227–33. <https://doi.org/10.1055/s-0038-1641200>.
- Levy, David H., Robert H. Dolin, John E. Mattison, Kent A. Spackman, and Keith E. Campbell. 1998. “Computer-Facilitated Collaboration: Experiences Building SNOMED-RT.” *Proceedings of the AMIA Symposium*, 870–74.
- Liu, Fangyu, Ehsan Shareghi, Zaiqiao Meng, Marco Basaldella, and Nigel Collier. 2021. “Self-Alignment Pretraining for Biomedical Entity Representations.” *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT)*, 4228–38.
- Miranda-Escalada, Antonio, Luis Gascó, Salvador Lima-López, et al. 2022. “Overview of DisTEMIST at BioASQ: Automatic Detection and Normalization of Diseases from Clinical Texts.” *Working Notes of CLEF 2022 – Conference and Labs of the Evaluation Forum*.
- Ngo, Duy-Hoa, Madonna Kemp, Donna Truran, Bevan Koopman, and Ale-

- jandro Metke-Jimenez. 2021. “Semantic Search for Large Scale Clinical Ontologies.” *AMIA Annual Symposium Proceedings* 2021: 910–19.
- O’Neil, M., C. Payne, and J. Read. 1995. “Read Codes Version 3: A User Led Terminology.” *Methods of Information in Medicine* 34 (1-2): 187–92.
- Olier, Ivan, David A. Springate, Darren M. Ashcroft, et al. 2016. “Modelling Conditions and Health Care Processes in Electronic Health Records: An Application to Severe Mental Illness with the Clinical Practice Research Datalink.” *PLOS ONE* 11 (2): e0146715. <https://doi.org/10.1371/journal.pone.0146715>.
- PostgreSQL Global Development Group. 2026. *PostgreSQL 16 Documentation: Chapter 12. Full Text Search*. PostgreSQL Global Development Group. <https://www.postgresql.org/docs/16/textsearch.html>.
- Remy, François, Kris Demuyne, and Thomas Demeester. 2024. “BioLORD-2023: Semantic Textual Representations Fusing LLM and Clinical Knowledge Graph Insights.” *Journal of the American Medical Informatics Association* 31 (9): 1844–55. <https://doi.org/10.1093/jamia/ocae029>.
- Rosenbloom, S. Trent, Joshua C. Denny, Hua Xu, Nancy Lorenzi, William W. Stead, and Kevin B. Johnson. 2011. “Data from Clinical Notes: A Perspective on the Tension Between Structure and Flexible Documentation.” *Journal of the American Medical Informatics Association* 18 (2): 181–86. <https://doi.org/10.1136/jamia.2010.007237>.
- Rosenbloom, S. Trent, Randolph A. Miller, Kevin B. Johnson, Peter L. Elkin, and Steven H. Brown. 2006. “Interface Terminologies: Facilitating Direct Entry of Clinical Data into Electronic Health Record Systems.” *Journal of the American Medical Informatics Association* 13 (3): 277–88. <https://doi.org/10.1197/jamia.M1957>.
- SNOMED International. 2024. *Do i Need an External Interface Terminology to Implement SNOMED CT?* Implementation Fact Sheets.
- VERIFY authors. 2024. “Use of SNOMED CT in Large Language Models: Scoping Review.” *JMIR Medical Informatics*.
- Warren, Judith J., Jane Collins, Carol Sorrentino, and James R. Campbell. 1998. “Just-in-Time Coding of the Problem List in a Clinical Environment.” *Proceedings of the AMIA Symposium*, 280–84.